

# Mercedes-Benz Greener Manufacturing

July 29, 2020

## DESCRIPTION

Reduce the time a Mercedes-Benz spends on the test bench.

Problem Statement Scenario: Since the first automobile, the Benz Patent Motor Car in 1886, Mercedes-Benz has stood for important automotive innovations. These include the passenger safety cell with a crumple zone, the airbag, and intelligent assistance systems. Mercedes-Benz applies for nearly 2000 patents per year, making the brand the European leader among premium carmakers. Mercedes-Benz is the leader in the premium car industry. With a huge selection of features and options, customers can choose the customized Mercedes-Benz of their dreams.

To ensure the safety and reliability of every unique car configuration before they hit the road, the company's engineers have developed a robust testing system. As one of the world's biggest manufacturers of premium cars, safety and efficiency are paramount on Mercedes-Benz's production lines. However, optimizing the speed of their testing system for many possible feature combinations is complex and time-consuming without a powerful algorithmic approach.

You are required to reduce the time that cars spend on the test bench. Others will work with a dataset representing different permutations of features in a Mercedes-Benz car to predict the time it takes to pass testing. Optimal algorithms will contribute to faster testing, resulting in lower carbon dioxide emissions without reducing Mercedes-Benz's standards.

Following actions should be performed:

If for any column(s), the variance is equal to zero, then you need to remove those variable(s).

- Check for null and unique values for test and train sets.
- Apply label encoder.
- Perform dimensionality reduction.
- Predict your test\_df values using XGBoost.

```
[1]: # Importing library

import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn import preprocessing # Import Label Encoder

[2]: # Read csv
train_df = pd.read_csv('train.csv')
test_df = pd.read_csv('test.csv')
```

```
print(train_df.shape) # Find Number of rows and columns
print(train_df.columns)

print(test_df.shape) # Find Number of rows and columns
print(test_df.columns)

train_df.head() # Show first 5 records
```

```
(4209, 378)
Index(['ID', 'y', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8',
      ...
      'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
      'X385'],
      dtype='object', length=378)
(4209, 377)
Index(['ID', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8', 'X10',
      ...
      'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
      'X385'],
      dtype='object', length=377)
```

```
[2]:
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	\
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0
2	0	1	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

[5 rows x 378 columns]

```
[3]: # Describe the dataset i.r.t its data Distribution
```

```
train_df.describe()
```

```
[3]:
```

	ID	y	X10	X11	X12	\
count	4209.000000	4209.000000	4209.000000	4209.0	4209.000000	
mean	4205.960798	100.669318	0.013305	0.0	0.075077	
std	2437.608688	12.679381	0.114590	0.0	0.263547	
min	0.000000	72.110000	0.000000	0.0	0.000000	

25%	2095.000000	90.820000	0.000000	0.0	0.000000
50%	4220.000000	99.150000	0.000000	0.0	0.000000
75%	6314.000000	109.010000	0.000000	0.0	0.000000
max	8417.000000	265.320000	1.000000	0.0	1.000000

	X13	X14	X15	X16	X17	...	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	...	
mean	0.057971	0.428130	0.000475	0.002613	0.007603	...	
std	0.233716	0.494867	0.021796	0.051061	0.086872	...	
min	0.000000	0.000000	0.000000	0.000000	0.000000	...	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
75%	0.000000	1.000000	0.000000	0.000000	0.000000	...	
max	1.000000	1.000000	1.000000	1.000000	1.000000	...	

	X375	X376	X377	X378	X379	...	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	...	
mean	0.318841	0.057258	0.314802	0.020670	0.009503	...	
std	0.466082	0.232363	0.464492	0.142294	0.097033	...	
min	0.000000	0.000000	0.000000	0.000000	0.000000	...	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
75%	1.000000	0.000000	1.000000	0.000000	0.000000	...	
max	1.000000	1.000000	1.000000	1.000000	1.000000	...	

	X380	X382	X383	X384	X385	...	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	...	
mean	0.008078	0.007603	0.001663	0.000475	0.001426	...	
std	0.089524	0.086872	0.040752	0.021796	0.037734	...	
min	0.000000	0.000000	0.000000	0.000000	0.000000	...	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
75%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
max	1.000000	1.000000	1.000000	1.000000	1.000000	...	

[8 rows x 370 columns]

**0.0.1** If for any column(s), the variance is equal to zero, then you need to remove those variable(s).

```
[4]: # Check the variance
```

```
train_df.var()
```

... ..

```
(train_df.var() == 0)
```

...

```
(train_df.var() == 0).values
```

```
array([False, False, False,  True, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False,  True, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False,  True, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
```

```
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, True, False, True, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, True, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, True, True, False, False,
True, False, False, False, True, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
True, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, True,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False])
```

```
[7]: variance_with_zero = train_df.var()[train_df.var()==0].index.values
variance_with_zero
```

```
[7]: array(['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290',
'X293', 'X297', 'X330', 'X347'], dtype=object)
```

```
[8]: # Drop zero variance variables

train_df = train_df.drop(variance_with_zero, axis=1)
```

```
[9]: print(train_df.shape)
```

```
(4209, 366)
```

```
[10]: # As ID column is irrelevant for our prediction hence we drop this column

train_df = train_df.drop(['ID'], axis=1)
```

```
[11]: train_df.head()
```



```
X380    False
X382    False
X383    False
X384    False
X385    False
Length: 365, dtype: bool
```

```
[14]: test_df.isnull().sum().values
```

[illegible]

```
[15]: # Find unique records
```

```
train_df.nunique()
```

```
[15]: y      2545
      X0      47
      X1      27
      X2      44
      X3       7
      ...
      X380     2
      X382     2
      X383     2
      X384     2
      X385     2
      Length: 365, dtype: int64
```

### 0.0.3 Filter out the columns having object datatype

```
[16]: object_datatypes = train_df.select_dtypes(include=[object])
      object_datatypes
```

```
[16]:
```

	X0	X1	X2	X3	X4	X5	X6	X8
0	k	v	at	a	d	u	j	o
1	k	t	av	e	d	y	l	o
2	az	w	n	c	d	x	j	x
3	az	t	n	f	d	x	l	e
4	az	v	n	f	d	h	d	n
...	..	..	..	..	..	..	..	..
4204	ak	s	as	c	d	aa	d	q
4205	j	o	t	d	d	aa	h	h
4206	ak	v	r	a	d	aa	g	e
4207	al	r	e	f	d	aa	l	u
4208	z	r	ae	c	d	aa	g	w

[4209 rows x 8 columns]

```
[17]: object_datatype_columns = object_datatypes.columns
      object_datatype_columns
```

```
[17]: Index(['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8'], dtype='object')
```

### 0.0.4 Apply label encoder.

```
[18]: # Initialize Label Encoder object

label_encoder = preprocessing.LabelEncoder()
train_df['X0'].unique()
```

```
[18]: array(['k', 'az', 't', 'al', 'o', 'w', 'j', 'h', 's', 'n', 'ay', 'f', 'x',
        'y', 'aj', 'ak', 'am', 'z', 'q', 'at', 'ap', 'v', 'af', 'a', 'e',
        'ai', 'd', 'aq', 'c', 'aa', 'ba', 'as', 'i', 'r', 'b', 'ax', 'bc',
        'u', 'ad', 'au', 'm', 'l', 'aw', 'ao', 'ac', 'g', 'ab'],
        dtype=object)
```

```
[19]: # Encode and transform object data to interger

train_df['X0'] = label_encoder.fit_transform(train_df['X0'])
```

```
[20]: train_df['X0'].unique()
```



```
[20]: array([32, 20, 40,  9, 36, 43, 31, 29, 39, 35, 19, 27, 44, 45,  7,  8, 10,
          46, 37, 15, 12, 42,  5,  0, 26,  6, 25, 13, 24,  1, 22, 14, 30, 38,
          21, 18, 23, 41,  4, 16, 34, 33, 17, 11,  3, 28,  2])
```

```
[21]: # Apply same for all columns having object type data
```

```
train_df['X1'] = label_encoder.fit_transform(train_df['X1'])
train_df['X2'] = label_encoder.fit_transform(train_df['X2'])
train_df['X3'] = label_encoder.fit_transform(train_df['X3'])
train_df['X4'] = label_encoder.fit_transform(train_df['X4'])
train_df['X5'] = label_encoder.fit_transform(train_df['X5'])
train_df['X6'] = label_encoder.fit_transform(train_df['X6'])
train_df['X8'] = label_encoder.fit_transform(train_df['X8'])
```

```
[22]: train_df.head()
```

```
[22]:
```

	y	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	\
0	130.81	32	23	17	0	3	24	9	14	0	...	0	0	1	0	
1	88.53	32	21	19	4	3	28	11	14	0	...	1	0	0	0	
2	76.26	20	24	34	2	3	27	9	23	0	...	0	0	0	0	
3	80.62	20	21	34	5	3	27	11	4	0	...	0	0	0	0	
4	78.02	20	23	34	5	3	12	3	13	0	...	0	0	0	0	

	X379	X380	X382	X383	X384	X385
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	1	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0

[5 rows x 365 columns]

### 0.0.5 Perform dimensionality reduction (PCA)

```
[23]: from sklearn.decomposition import PCA
```

```
[24]: # PCA with 95%
```

```
sklearn_pca = PCA(n_components=0.95)
```

```
[25]: sklearn_pca.fit(train_df)
```

```
[25]: PCA(copy=True, iterated_power='auto', n_components=0.95, random_state=None,
        svd_solver='auto', tol=0.0, whiten=False)
```

```
[26]: x_train_transformed = sklearn_pca.transform(train_df)
```

```
[27]: print(x_train_transformed.shape)
```

```
(4209, 6)
```

```
[28]: # PCA with 98%
```

```
sklearn_pca_98 = PCA(n_components=0.98)
```

```
[29]: sklearn_pca_98.fit(train_df)
```

```
[29]: PCA(copy=True, iterated_power='auto', n_components=0.98, random_state=None,  
      svd_solver='auto', tol=0.0, whiten=False)
```

```
[30]: x_train_transformed_98 = sklearn_pca_98.transform(train_df)  
      print(x_train_transformed_98.shape)
```

```
(4209, 12)
```

```
[31]: train_df.y
```

```
[31]: 0      130.81  
      1      88.53  
      2      76.26  
      3      80.62  
      4      78.02  
      ...  
      4204    107.39  
      4205    108.77  
      4206    109.22  
      4207     87.48  
      4208    110.85  
      Name: y, Length: 4209, dtype: float64
```

## 0.0.6 Train and Test split on Train dataset

```
[32]: X = train_df.drop('y', axis=1)  
      y = train_df.y  
      xtrain,xtest,ytrain,ytest = train_test_split(X,y,test_size=0.3,random_state=42)
```

```
[33]: print(xtrain)  
      print(xtrain.shape)
```

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	...	X375	X376	X377	X378	\
370	35	13	16	1	3	9	6	19	0	0	...	0	0	0	0	
3392	15	10	16	2	3	23	9	16	0	0	...	0	0	1	0	
2208	31	3	16	2	3	15	2	21	0	0	...	0	0	1	0	

3942	35	20	8	6	3	26	6	14	0	1	...	1	0	0	0
1105	36	13	16	5	3	1	6	0	0	0	...	0	0	0	0
...	..	..	..	..	..	..	..	..	...	...	...	...	...	...	...
3444	31	10	16	2	3	22	11	17	0	0	...	0	0	1	0
466	20	25	25	2	3	9	9	9	0	0	...	0	0	0	0
3092	45	24	3	2	3	21	8	2	0	0	...	1	0	0	0
3772	45	19	8	5	3	25	8	1	0	1	...	0	0	0	0
860	22	1	7	2	3	5	9	17	0	0	...	1	0	0	0

	X379	X380	X382	X383	X384	X385
370	0	0	0	0	0	0
3392	0	0	0	0	0	0
2208	0	0	0	0	0	0
3942	0	0	0	0	0	0
1105	0	0	0	0	0	0
...	...	...	...	...	...	...
3444	0	0	0	0	0	0
466	0	0	1	0	0	0
3092	0	0	0	0	0	0
3772	0	0	0	0	0	0
860	0	0	0	0	0	0

[2946 rows x 364 columns]  
(2946, 364)

```
[34]: print(ytrain)
      print(ytrain.shape)
```

```
370      95.13
3392     117.36
2208     109.01
3942      93.77
1105     103.41
...
3444     109.42
466       78.25
3092      92.18
3772      91.92
860       87.71
Name: y, Length: 2946, dtype: float64
(2946,)
```

```
[35]: print(xtest)
      print(xtest.shape)
```

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	...	X375	X376	X377	X378	\
1073	9	16	7	5	3	6	9	11	0	0	...	0	0	0	0	
144	27	13	3	5	3	13	8	22	0	0	...	0	0	0	0	

2380	31	1	21	2	3	18	11	14	1	0	...	1	0	0	0
184	20	25	22	2	3	13	9	11	0	0	...	0	0	0	0
2587	8	23	8	3	3	17	8	17	0	0	...	0	0	0	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
2493	27	20	16	2	3	18	10	5	0	0	...	0	0	1	0
3388	40	19	24	5	3	23	3	19	0	0	...	0	0	0	0
3997	22	3	7	0	3	26	6	18	0	0	...	0	0	1	0
383	40	1	16	6	3	9	8	0	0	0	...	1	0	0	0
3364	27	4	33	2	3	23	6	24	0	0	...	0	0	1	0

	X379	X380	X382	X383	X384	X385
1073	0	0	0	0	0	0
144	0	0	0	0	0	0
2380	0	0	0	0	0	0
184	0	0	1	0	0	0
2587	0	0	0	0	0	0
...	...	...	...	...	...	...
2493	0	0	0	0	0	0
3388	0	0	0	0	0	0
3997	0	0	0	0	0	0
383	0	0	0	0	0	0
3364	0	0	0	0	0	0

[1263 rows x 364 columns]  
(1263, 364)

```
[36]: # PCA with 95% for xtrain
```

```
pca_xtrain = PCA(n_components=0.95)
pca_xtrain.fit(xtrain)
```

```
[36]: PCA(copy=True, iterated_power='auto', n_components=0.95, random_state=None,
      svd_solver='auto', tol=0.0, whiten=False)
```

```
[37]: pca_xtrain_transformed = pca_xtrain.transform(xtrain)
      print(pca_xtrain_transformed.shape)
```

(2946, 6)

```
[38]: # PCA with 95% for xtest
```

```
pca_xtest = PCA(n_components=0.95)
pca_xtest.fit(xtest)
```

```
[38]: PCA(copy=True, iterated_power='auto', n_components=0.95, random_state=None,
      svd_solver='auto', tol=0.0, whiten=False)
```

```
[39]: pca_xtest_transformed = pca_xtest.transform(xtest)
print(pca_xtest_transformed.shape)
```

(1263, 6)

```
[40]: print(pca_xtest.explained_variance_)
print(pca_xtest.explained_variance_ratio_)
```

```
[206.79524961 120.24273955  67.64680756  61.94375666  48.08214872
  8.7271811 ]
[0.38517942 0.22396563 0.12599979 0.11537722 0.08955841 0.01625536]
```

### 0.0.7 PCA for test\_df dataset

```
[41]: test_df
```

```
[41]:
```

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	\
0	1	az	v	n	f	d	t	a	w	0	...	0	0	0	1	
1	2	t	b	ai	a	d	b	g	y	0	...	0	0	1	0	
2	3	az	v	as	f	d	a	j	j	0	...	0	0	0	1	
3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
4204	8410	aj	h	as	f	d	aa	j	e	0	...	0	0	0	0	
4205	8411	t	aa	ai	d	d	aa	j	y	0	...	0	1	0	0	
4206	8413	y	v	as	f	d	aa	d	w	0	...	0	0	0	0	
4207	8414	ak	v	as	a	d	aa	c	q	0	...	0	0	1	0	
4208	8416	t	aa	ai	c	d	aa	g	r	0	...	1	0	0	0	
		X379	X380	X382	X383	X384	X385									
0		0	0	0	0	0	0									
1		0	0	0	0	0	0									
2		0	0	0	0	0	0									
3		0	0	0	0	0	0									
4		0	0	0	0	0	0									
...	...	...	...	...	...											
4204		0	0	0	0	0	0									
4205		0	0	0	0	0	0									
4206		0	0	0	0	0	0									
4207		0	0	0	0	0	0									
4208		0	0	0	0	0	0									

[4209 rows x 377 columns]

```
[42]: test_object_datatypes = test_df.select_dtypes(include=[object])
test_object_datatypes
```

```
[42]:      X0  X1  X2 X3 X4  X5 X6 X8
0      az  v   n   f   d   t   a   w
1      t   b  ai   a   d   b   g   y
2      az  v   as  f   d   a   j   j
3      az  l   n   f   d   z   l   n
4      w   s  as  c   d   y   i   m
...
4204   aj   h  as  f   d  aa   j   e
4205   t   aa ai   d   d  aa   j   y
4206   y   v  as  f   d  aa   d   w
4207  ak   v  as  a   d  aa   c   q
4208   t   aa ai   c   d  aa   g   r
```

[4209 rows x 8 columns]

```
[43]: test_df['X0'] = label_encoder.fit_transform(test_df['X0'])
test_df['X1'] = label_encoder.fit_transform(test_df['X1'])
test_df['X2'] = label_encoder.fit_transform(test_df['X2'])
test_df['X3'] = label_encoder.fit_transform(test_df['X3'])
test_df['X4'] = label_encoder.fit_transform(test_df['X4'])
test_df['X5'] = label_encoder.fit_transform(test_df['X5'])
test_df['X6'] = label_encoder.fit_transform(test_df['X6'])
test_df['X8'] = label_encoder.fit_transform(test_df['X8'])
```

```
[44]: print(test_df)
print(test_df.shape)
```

```
      ID  X0  X1  X2  X3  X4  X5  X6  X8  X10  ...  X375  X376  X377  X378  \
0      1  21  23  34   5   3  26   0  22   0  ...    0    0    0    1
1      2  42   3   8   0   3   9   6  24   0  ...    0    0    1    0
2      3  21  23  17   5   3   0   9   9   0  ...    0    0    0    1
3      4  21  13  34   5   3  31  11  13   0  ...    0    0    0    1
4      5  45  20  17   2   3  30   8  12   0  ...    1    0    0    0
...
4204  8410   6   9  17   5   3   1   9   4   0  ...    0    0    0    0
4205  8411  42   1   8   3   3   1   9  24   0  ...    0    1    0    0
4206  8413  47  23  17   5   3   1   3  22   0  ...    0    0    0    0
4207  8414   7  23  17   0   3   1   2  16   0  ...    0    0    1    0
4208  8416  42   1   8   2   3   1   6  17   0  ...    1    0    0    0
```

```
      X379  X380  X382  X383  X384  X385
0          0     0     0     0     0     0
1          0     0     0     0     0     0
2          0     0     0     0     0     0
3          0     0     0     0     0     0
4          0     0     0     0     0     0
...      ...     ...     ...     ...     ...
```

```

4204    0    0    0    0    0    0
4205    0    0    0    0    0    0
4206    0    0    0    0    0    0
4207    0    0    0    0    0    0
4208    0    0    0    0    0    0

```

```

[4209 rows x 377 columns]
(4209, 377)

```

```
[45]: test_df = test_df.drop('ID',axis=1)
```

```
[46]: # PCA with 95% for test_df

pca_test_df = PCA(n_components=0.95)
pca_test_df.fit(test_df)
```

```
[46]: PCA(copy=True, iterated_power='auto', n_components=0.95, random_state=None,
      svd_solver='auto', tol=0.0, whiten=False)
```

```
[47]: pca_test_df_transformed = pca_test_df.transform(test_df)
      print(pca_test_df_transformed.shape)
```

```
(4209, 6)
```

```
[48]: print(pca_test_df.explained_variance_)
      print(pca_test_df.explained_variance_ratio_)
```

```

[247.07875325 100.33535335  77.48364816  62.33258307  48.95689653
   8.14203723]
[0.43515102 0.17670897 0.13646292 0.10977912 0.08622208 0.01433962]

```

```
[49]: y
```

```

[49]: 0      130.81
      1      88.53
      2      76.26
      3      80.62
      4      78.02
      ...
      4204    107.39
      4205    108.77
      4206    109.22
      4207     87.48
      4208    110.85
      Name: y, Length: 4209, dtype: float64

```

### 0.0.8 Perform XGboost

```
[78]: from sklearn import svm
      from sklearn import model_selection
      import xgboost as xgb
```

```
[80]: model = xgb.XGBRegressor(objective="reg:linear",learning_rate=0.1)
      model.fit(pca_xtrain, ytrain) # I am getting a small error here, unable to
      ↪ solve.Please help me with solution.
      y_pred = model.predict(pca_x_test)
      y_pred
      model.predict(pca_test_df)
```

```
[ ]: ** End **
```